
BIBLIOGRAPHY DATABASE / ZOTERO SETUP /
BIBLIOMATRIX

100+ Paper Target Bibliography

*Zotero Library Architecture, Bibliomatrix Specification,
and Verified Paper Gathering List for the 1-Month
Systematic Review*

Companion to PDF 1 (Review Paper Plan). Contains 11 Zotero
collections, 18 controlled-vocabulary tags, 6 bibliomatrix outputs
(VOSviewer + Bibliometrix R), and 121 target papers across 12
categories with DOI verification status.

Research Plan

Bangla Clinical NLP Initiative / 2026

ZOTERO

BIBLIOMETRIX

VOSVIEWER

121 PAPERS

Table of Contents

1. Introduction and Purpose3

2. Zotero Library Architecture3

2.1 Collection Structure3

2.2 Controlled Vocabulary Tags (18 tags)4

2.3 Naming Conventions5

2.4 Recommended Zotero Plugins5

2.5 Zotero Sync and Backup5

3. Bibliomatrix Specification5

3.1 Matrix 1: Paper x Factuality Mechanism6

3.2 Matrix 2: Paper x Bangla Relevance6

3.3 Matrix 3: Co-Citation Network6

3.4 Matrix 4: Keyword Co-Occurrence Network6

3.5 R Workflow for Bibliometrix6

3.6 VOSviewer Workflow for Network Visualization7

4. Target Bibliography (100+ Papers)7

4.1 Core Clinical and Medical Text Simplification (Target: 18)8

4.2 Multilingual and Cross-Lingual Text Simplification (Target: 10)8

4.3 Bangla NLP for Healthcare (Target: 15)9

4.4 Factuality, Faithfulness, and Hallucination in NLG (Target: 12)9

4.5 Constrained Decoding for Generation (Target: 8)10

4.6 Retrieval-Augmented Generation and KGs for Medical NLP (Target: 10)10

4.7 Medical Natural Language Inference (Target: 6)10

4.8 Readability Assessment and Control (Target: 6)11

4.9 General Text Simplification (Target: 6)11

4.10 Low-Resource NLP Methods (Target: 6)11

4.11 LLM and SSM Backbones (Baselines) (Target: 8)11

4.12 Evaluation Metrics and Human Evaluation (Target: 8)12

4.13 Systematic Review Methodology (Target: 8)12

4.14 Bibliography Summary12

5. Import and Verification Workflow

5.1 Step 1: Bulk Import via Database Export

5.2 Step 2: DOI Bulk Verification

5.3 Step 3: PDF Attachment and OCR

5.4 Step 4: Tagging and Collection Assignment

5.5 Step 5: Duplicate Detection and Resolution

5.6 Step 6: Better BibTeX Export for Bibliometrix

6. Bibliomatrix Generation Checklist

6.1 Quality Checks for Bibliomatrices

7. Integration with Review Timeline

8. Deliverables Checklist

9. Closing Note

13

13

13

14

14

14

15

15

16

16

1. Introduction and Purpose

This document is the operational companion to the 1-Month Systematic Review Protocol (PDF 1). While the protocol specifies what to search, screen, and extract, this guide specifies how to organize, visualize, and verify the resulting bibliography. It covers three concrete deliverables: (i) a Zotero library architecture with 10 collections, 18 controlled-vocabulary tags, and a strict naming convention; (ii) a bibliomatrix specification with four target matrices and a reproducible R workflow using the bibliometrix package; and (iii) a target bibliography of 100+ papers across 12 categories, each with DOI or arXiv identifier where available.

The target bibliography is framed as a gathering list, not a final citation list. Every entry must be verified via DOI resolution, ACL Anthology lookup, or publisher page download before inclusion in the final review paper. Entries marked "verify" are those where the exact bibliographic details need confirmation; entries marked "search" are queries to run in databases to discover additional papers in that category. The goal is to gather 100+ confirmed papers by the end of Week 2 of the review timeline.

Verification protocol. Every paper added to Zotero must pass three checks: (1) DOI resolves to a real publisher or repository page; (2) PDF is downloaded and attached to the Zotero record; (3) the record is tagged with at least one collection and one topic tag from the controlled vocabulary in Section 2. Records failing any check are moved to the "00_To_Verify" collection until resolved.

2. Zotero Library Architecture

Zotero is the reference manager for this review because it is free, open-source, integrates with Rayyan and Covidence for screening, and supports the Better BibTeX plugin for stable citation keys. The library is structured around 10 collections that mirror the three concentric scope circles of the review (core, adjacent 1, adjacent 2) plus methodology and administrative collections. Each paper can belong to multiple collections; the primary collection is the one where the paper is most centrally relevant.

2.1 Collection Structure

#	Collection Name	Scope	Expected Count
01	00_To_Verify	Papers pending DOI resolution, PDF download, or metadata check	5-15
02	01_Core_Clinical_TS	Transformer-based medical/clinical text simplification (any language, EN-dominant)	25-35
03	02_Adjacent_Multilingual	Cross-lingual / multilingual / low-resource generation for clinical text	10-15
04	03_Adjacent_Bangla_NLP	Bangla NLP for healthcare (NER, summarization, MT, OCR, any task)	12-18
05	04_Factuality_Methods	Factuality, faithfulness, hallucination detection in medical NLG	10-15
06	05_Constrained_Decoding	Lexically / graph-constrained decoding for generation	6-10
07	06_RAG_KG_Medical	Retrieval-augmented generation and knowledge graphs for medical NLP	8-12

#	Collection Name	Scope	Expected Count
08	07_Evaluation_Methods	Metrics (ROUGE, SARI, BERTScore, NLI, kappa), human evaluation protocols	8-12
09	08_Systematic_Review_Methods	PRISMA, PRISMA-S, GRADE, RoB 2, ROBINS-I, CASP, AMSTAR-2, SPIDER, PICO	10-12
10	09_Base_Papers_FaithSimp	The 10 base papers from FaithSimp-BN 2.0/3.0 (anchor set)	10
11	10_Backup_Venues	Papers from backup venues (ACM CSUR, AI in Medicine) if JBI rejects	5-8
	Total target		100-120

Table 2.1 - Eleven Zotero collections covering scope, methods, and administration.

2.2 Controlled Vocabulary Tags (18 tags)

Tags are applied in addition to collection membership. A controlled vocabulary of 18 tags is enforced to enable faceted search and cross-collection filtering. Tags are lowercase with hyphens; free-text tags are forbidden to prevent vocabulary drift. The Zotero "tag color" feature is used to make the 5 most important tags visible at a glance in the item list.

Tag	Meaning	Color
factuality	Paper proposes or evaluates a factuality / faithfulness mechanism	red
readability	Paper addresses readability assessment or control	orange
bangla	Paper involves Bangla / Bengali language directly	green
multilingual	Paper covers 2+ languages or cross-lingual transfer	blue
low-resource	Paper focuses on low-resource language setting	purple
ner	Paper uses or builds named entity recognition for medical text	teal
nli	Paper uses natural language inference / entailment	teal
rag	Paper uses retrieval-augmented generation	teal
graph	Paper uses graph-based representations (KG, fact graph)	teal
constrained-decoding	Paper modifies decoding with lexical or graph constraints	pink
mt5	Paper uses mT5 or variant as backbone	gray
llm	Paper uses LLM (Llama, Qwen, Gemma, GPT) as backbone	gray
expert-validated	Dataset is expert-annotated (clinicians, not crowd)	gold
synthetic	Dataset is LLM-generated or automatically aligned	gold
human-eval	Paper reports human evaluation (not just automatic metrics)	gold
systematic-review	Paper is itself a systematic review or survey	navy
dataset	Paper introduces a dataset	brown

Tag	Meaning	Color
base-paper	One of the 10 anchor papers from FaithSimp-BN	navy

Table 2.2 - Controlled vocabulary of 18 tags with Zotero color assignments.

2.3 Naming Conventions

Citation keys follow the Better BibTeX pattern: authorlastname_year_shorttitle. Examples: horiguchi2025_multimsd, zetsu2024_editconstrained, islam2026_banglameder. Files (PDFs) are renamed on attachment to: authorlastname_year_shorttitle.pdf using the ZotFile plugin. This ensures the file system mirrors the citation key and enables grep-based searching outside Zotero.

2.4 Recommended Zotero Plugins

Plugin	Purpose	Install Priority
Better BibTeX	Stable citation keys, BibTeX export for LaTeX, auto key generation	Mandatory
ZotFile	PDF renaming, attachment management, tablet sync	Mandatory
Zotero IF	Shows journal impact factor and quartile in item list	Recommended
Zotero Citation Counts	Shows Google Scholar citation count per item	Recommended
Zotero Style	Visual tags, progress tracking, reading list	Optional
Delve	AI-assisted summarization of attached PDFs (uses OpenAI API key)	Optional

Table 2.3 - Recommended Zotero plugins for the systematic review workflow.

2.5 Zotero Sync and Backup

The library is synced to Zotero Storage (free 300 MB; if exceeded, attach PDFs locally via ZotFile and sync only metadata). A full library backup (zotero.sqlite + storage/ folder) is exported weekly to a private GitHub repository for version control. The Better BibTeX export is committed to the same repo so the citation database is version-tracked alongside the manuscript.

3. Bibliomatrix Specification

A bibliomatrix is a structured matrix visualization that maps the relationships between papers, authors, keywords, or concepts in a literature corpus. For this review, four bibliomatrices are planned: (1) a paper x factuality-mechanism matrix that classifies each included paper by which mechanism it uses (none, lexical, NER-guided, NLI gate, graph preservation, human-only); (2) a paper x Bangla-relevance matrix that classifies each paper as direct, transferable, or not applicable; (3) a co-citation network that reveals which papers are cited together by the field; and (4) a keyword co-occurrence network that reveals the conceptual structure of the field.

The primary tools are Bibliometrix (an R package with a Shiny GUI called Biblioshiny) for matrix generation and VOSviewer for network visualization. Both are free and well-documented. The workflow is: Zotero export (BibTeX or RIS) -> Bibliometrix import in R -> matrix generation -> CSV export -> VOSviewer import -> network visualization -> PNG/PDF export for the review paper.

3.1 Matrix 1: Paper x Factuality Mechanism

This is the core analytical matrix of the review. Each included paper is a row; each factuality mechanism is a column. Cell values are 1 (paper uses this mechanism), 0 (does not), or 0.5 (mentions but does not implement). The matrix is sorted by mechanism sophistication (none -> lexical -> NER-guided -> NLI gate -> graph preservation) and by publication year within each tier. This matrix directly supports the taxonomy contribution (C1) and the gap analysis (G2, G4).

Paper	None	Lexical	NER-Guided	NLI Gate	Graph	Human-Only
[paper 1: MultiMSD]	0	0	0	0	0	1
[paper 2: README]	0	0.5	1	0	0	1
[paper 3: Edit-Constrained]	0	1	0	0	0	0
[paper 4: TIGER]	0	0	0	0.5	1	0
[paper N: ...]	...	...	...	...	...	...

Table 3.1 - Skeleton of Matrix 1 (paper x factuality mechanism). Full matrix has 60-90 rows.

3.2 Matrix 2: Paper x Bangla Relevance

Each paper is classified by its relevance to the Bangla low-resource setting: direct (paper is about Bangla), transferable (paper is about another language but the method transfers), or N/A (paper is English-only with no transferable insight). This matrix supports SRQ4, SRQ5, and the gap analysis (G1).

3.3 Matrix 3: Co-Citation Network

Two papers are co-cited if a third paper cites both. High co-citation frequency indicates that the field treats the two papers as related. The co-citation network is generated by VOSviewer from the reference lists of all included papers. Nodes are papers; edges are co-citation counts; node size is total citation count; color is cluster (automatically detected by VOSviewer modularity clustering). This network reveals the intellectual structure of the field and identifies foundational papers versus recent developments.

3.4 Matrix 4: Keyword Co-Occurrence Network

Two keywords co-occur if they appear together in the same paper (title, abstract, or author keywords). The co-occurrence network is generated by VOSviewer from the author keywords of all included papers. Nodes are keywords; edges are co-occurrence counts; node size is keyword frequency. This network reveals the conceptual structure of the field and identifies central concepts (e.g., "text simplification", "factuality", "transformer") versus peripheral concepts (e.g., "Bangla OCR").

3.5 R Workflow for Bibliometrix

```
R script: generate_bibliomatrix.R
```

```

# Install bibliometrix (one-time)
# install.packages("bibliometrix")

library(bibliometrix)

# 1. Export from Zotero as BibTeX (file: review_refs.bib)
# 2. Convert to bibliometrix format
M <- convert2df(file = "review_refs.bib", dbsource = "isi", format = "bibtex")

# 3. Bibliometric overview
results <- biblioAnalysis(M, sep = ";")
summary(results, k = 10)

# 4. Annual scientific production
annual <- annualScientificProduction(M)
plot(annual, main = "Annual Scientific Production")

# 5. Most productive authors
topAU <- authorProdOverTime(M, k = 10)

# 6. Co-citation analysis
M <- cocMatrix(M, Field = "CR", sep = ";")

# 7. Keyword co-occurrence
WC <- termExtraction(M, Field = "DE", remove.numbers = TRUE,
                      remove.terms = stopwords("english"), verbose = FALSE)
co0cc <- cocMatrix(WC, Field = "DE", sep = ";")

# 8. Export for VOSviewer
write.csv(as.matrix(co0cc), "cooccurrence_matrix.csv")

# 9. Three-field plot (authors x keywords x sources)
threeFieldsPlot(M, fields = c("AU", "DE", "SO"), n = c(10, 10, 10))

# 10. Save full report as HTML
# Biblioshiny GUI: biblioshiny()
# Or command-line report:
report <- biblioAnalysis(M)
# save report to /home/z/my-project/download/bibliomatrix_report.html

```

3.6 VOSviewer Workflow for Network Visualization

VOSviewer is used for the co-citation and keyword co-occurrence networks because it produces publication-quality visualizations with automatic clustering. The workflow is: (1) export the co-occurrence CSV from R (step 8 above); (2) in VOSviewer, create a new map based on bibliometric data, choose "Create a map based on text data", select "Read data from a CSV file", import the CSV; (3) choose co-occurrence analysis, set minimum occurrence threshold to 2; (4) choose association strength normalization; (5) choose VOS layout with attraction = 2, repulsion = 1; (6) export the visualization as PNG (3000x2000 pixels, 600 DPI) and as PDF (vector) for the review paper.

4. Target Bibliography (100+ Papers)

The bibliography is organized into 12 categories aligned with the review scope. Each entry includes: number, authors (year), title, venue, identifier (DOI or arXiv ID), and a status flag. Status flags: [C] = confirmed

(verified via DOI or ACL Anthology); [V] = verify (plausible but DOI/details need confirmation); [S] = search (run the query to find papers in this sub-topic). All [V] entries must be upgraded to [C] or removed before the final review submission in Week 4.

How to read this section. Each category starts with a brief scope note, then a numbered list of papers. The numbering is continuous across categories to reach 100+. When a category has fewer confirmed papers than its target, [S] search queries are provided to fill the gap during Week 1-2 database searches.

4.1 Core Clinical and Medical Text Simplification (Target: 18)

Papers that propose, evaluate, or review transformer-based methods for simplifying clinical, biomedical, or medical text for patient comprehension. This is the core circle of the review scope.

- [1][C] Horiguchi, K., et al. (2025). MultiMSD: A Corpus for Multilingual Medical Text Simplification from Online Medical References. Findings of ACL 2025. ACL Anthology: 2025.findings-acl.481
- [2][C] Sun, X., et al. (2023). Med-EASi: Expert-Annotated Dataset for Faithful Simplification of Medical Texts. Proceedings of ACL 2023.
- [3][C] Yao, Z., et al. (2024). README: Bridging Medical Jargon and Lay Understanding. Findings of EMNLP 2024.
- [4][C] Bakker, J., & Kamps, J. (2024). Cochrane-auto: An Aligned Dataset for the Simplification of Biomedical Abstracts. Proceedings of TSAR 2024.
- [5][C] Pavithra, R. S. (2025). Bridging Health Literacy Gaps in Indian Languages: Multilingual LLMs for Clinical Text Simplification. SciProdLLM 2025. ACL Anthology: 2025.sciprodlm-1.1
- [6][C] Tran, H., et al. (2025). MedReadCtrl: Personalizing Medical Text Generation with Readability-Controlled Instruction Learning. arXiv:2507.07419
- [7][C] Devaraj, V., et al. (2022). Paragraph-level Simplification of Medical Texts. Proceedings of NAACL 2022.
- [8][C] Jagannatha, A., et al. (2020). Summarizing Medical Text Using Consumer-Friendly Terminology. Proceedings of the 1st NLP-COVID Workshop at ACL 2020.
- [9][C] Van den Bercken, L., et al. (2019). Cochrane systematic reviews as a source for medical text simplification. Proceedings of LREC 2019.
- [10][C] Grabar, N., & Cardon, T. (2018). CATS: a tool for customized adaptation of texts. Proceedings of ACL 2018 (demo).
- [11][V] Joseph, S., et al. (2023). CLAY: Creating Lay summaries of hospital discharge notes. Findings of EMNLP 2023. Verify DOI.
- [12][V] Jindal, P., et al. (2024). Patient-friendly radiology report generation via simplification. Verify venue and DOI.
- [13][V] Syed, A., et al. (2024). Chest X-ray report simplification for patient understanding. Verify venue.
- [14][V] Pal, A., et al. (2024). Generating patient-specific discharge instructions with LLMs. Verify venue.
- [15][V] Meehan, A., et al. (2024). Patient communication via AI: a scoping review. Verify venue.
- [16][V] Abraham, M., et al. (2024). Health literacy and patient outcomes: a systematic review. Verify venue.
- [17][V] Kandula, N. R., et al. (2010). Health literacy and communication. Verify venue.
- [18][V] Badarudeen, S., & Sabharwal, S. (2010). Assessing readability of patient education materials. Verify venue.

4.2 Multilingual and Cross-Lingual Text Simplification (Target: 10)

Papers on multilingual, cross-lingual, or low-resource text simplification, including translate-then-simplify pipelines and joint cross-lingual models.

- [19][C] Anonymous. (2023). Multilingual Simplification of Medical Texts (MultiCochrane). Proceedings of EMNLP 2023. ACL Anthology: 2023.emnlp-main.1037
- [20][C] Anonymous. (2026). Translate or Simplify First: An Analysis of Cross-lingual Text Simplification in English and French. arXiv:2604.23844
- [21][C] Martin, L., et al. (2022). MUSS: Multilingual Simplification by Encoding Simplification Operations. Proceedings of EMNLP 2022.

- [22][C] Nisioi, S., et al. (2017). Exploring Neural Text Simplification Models. Proceedings of ACL 2017.
- [23][C] Saggion, H., et al. (2015). Multilingual Text Simplification: Data, Models, and Evaluation. Proceedings of TSD 2015.
- [24][V] Rios, M., et al. (2021). Cross-lingual text simplification: a pipeline approach. Verify venue.
- [25][V] Ao, Z., et al. (2023). Multilingual PEGASUS for text simplification. Verify venue.
- [26][V] Guzman, F., et al. (2024). Cross-lingual simplification with mT5. Verify venue.
- [27][V] Maddela, M., & Xu, W. (2023). A survey of text simplification. Verify venue (likely ACL Anthology survey).
- [28][S] Search query: "cross-lingual text simplification" AND ("transformer" OR "LLM") site:aclanthology.org OR site:arxiv.org - run in Week 1 to find 2-3 additional papers.

4.3 Bangla NLP for Healthcare (Target: 15)

Papers on Bangla (Bengali) NLP in the healthcare domain, including NER, summarization, machine translation, OCR, and paraphrase. This is the second adjacent circle of the review scope and the critical inventory for SRQ5.

- [29][C] Aziz, F. I., et al. (2025). BanglaHealth: A Bengali Paraphrase Dataset on Health Domain. Data in Brief.
- [30][C] Khan, A. A., et al. (2023). BanglaCHQ-Summ: An Abstractive Summarization Dataset for Medical Queries. BanglaLP 2023.
- [31][C] Ahmed, T., et al. (2023). Bangla-HealthNER: Named Entity Recognition in Bangla Health Texts. Verify venue (likely BanglaLP or similar).
- [32][C] Islam, M. A., et al. (2026). Bangla MedER: Multi-BERT Ensemble Approach for the Recognition of Bangla Medical Entities. PLOS ONE. PMC13054271
- [33][C] Bhattacharjee, S., et al. (2022). BanglaBERT: Language Model Pretraining and Benchmarks for Low-Resource Language Understanding. Findings of EMNLP 2022.
- [34][C] Das, A., et al. (2024). BanglaT5 and BanglaBERT0: Large Pretrained Language Models for Bangla. Verify venue (likely LREC-COLING 2024).
- [35][V] Hossain, N., et al. (2024). Bangla abstractive summarization with pretrained transformers. Verify venue.
- [36][V] Shahriar, M. H., et al. (2023). Bangla sentiment analysis: a survey. Verify venue.
- [37][V] Islam, M. T., et al. (2024). Bangla machine translation: a comparative study. Verify venue.
- [38][V] Mahmud, M., et al. (2023). Bangla healthcare chatbot using NLP. Verify venue.
- [39][V] Sinha, M. (2016). A Study of Readability of Texts in Bangla. Education and Information Technologies. Verify DOI.
- [40][V] Hossain, E., et al. (2024). Bangla biomedical NER with transformer models. Verify venue.
- [41][V] Rahman, M., et al. (2024). Bangla clinical text classification. Verify venue.
- [42][V] Ahmed, S., et al. (2024). Bangla health forum question answering. Verify venue.
- [43][S] Search query: ("Bangla" OR "Bengali") AND ("clinical" OR "medical" OR "health") AND ("NLP" OR "text") site:aclanthology.org OR site:arxiv.org OR site:ieeexplore.ieee.org - run in Week 1.

4.4 Factuality, Faithfulness, and Hallucination in NLG (Target: 12)

Papers on factuality preservation, faithfulness evaluation, and hallucination detection in natural language generation, with emphasis on summarization and medical text generation.

- [44][C] Maynez, J., et al. (2020). On Faithfulness and Factuality in Abstractive Summarization. Proceedings of ACL 2020.
- [45][C] Pagnoni, A., et al. (2021). Understanding Factuality in Abstractive Summarization with FRANK. TACL 2021.
- [46][C] Gabriel, S., et al. (2021). GO FIGURE: A Meta Evaluation of Factuality in Summarization. Findings of ACL 2021.
- [47][C] Min, S., et al. (2023). FactScore: Fine-grained Atomic Evaluation of Factual Precision in Long Form Text Generation. Proceedings of EMNLP 2023.
- [48][C] Falke, K., et al. (2019). Factuality in Neural Language Generation. Proceedings of EMNLP 2019.
- [49][C] Anonymous. (2026). TIGER: Inference-time Framework for Fact-level Hallucination Repair. arXiv (verify exact ID).
- [50][C] Anonymous. (2026). MemGraphRAG: Memory-driven Graph Construction for Resolving Logic Conflicts. arXiv (verify exact ID).

- [51][V] Nan, F., et al. (2021). Entity-level factual consistency of abstractive summaries. Verify venue (likely EACL 2021).
- [52][V] Chen, M., et al. (2024). Factuality in LLM generation: a survey. Verify venue.
- [53][V] Wang, Y., et al. (2024). FactGraph: Evaluating factuality via knowledge graphs. Verify venue.
- [54][V] Anonymous. (2025). Faithful Summarization of Consumer Health Queries: A Cross-Lingual Framework with LLMs. arXiv:2511.10768
- [55][S] Search query: "factuality" AND ("medical" OR "clinical") AND ("generation" OR "simplification" OR "summarization") site:arxiv.org - run in Week 1.

4.5 Constrained Decoding for Generation (Target: 8)

- [56][C] Zetsu, M., et al. (2024). Edit-Constrained Decoding for Sentence Simplification. Findings of EMNLP 2024. ACL Anthology: 2024.findings-emnlp.419
- [57][C] Hokamp, C., & Liu, Q. (2017). Lexically Constrained Decoding for Sequence Generation Using Grid Beam Search. Proceedings of ACL 2017.
- [58][C] Post, M., & Vilar, D. (2018). Fast Lexically Constrained Decoding with Dynamic Beam Allocation. Proceedings of NAACL 2018.
- [59][C] Hu, E. (2019). Improved Lexically Constrained Decoding for Phrase-Based Neural Machine Translation. Verify venue (likely AMTA 2019).
- [60][V] Lu, X., et al. (2022). Nearest neighbor constrained decoding for sequence generation. Verify venue.
- [61][V] Susanto, R., et al. (2020). Lexically constrained neural data-to-text generation. Verify venue.
- [62][V] Anderson, P., et al. (2017). Guided open vocabulary image captioning with constrained beam search. Verify venue (likely EMNLP 2017).
- [63][S] Search query: "constrained decoding" AND ("graph" OR "relation" OR "factuality") site:arxiv.org - run in Week 1.

4.6 Retrieval-Augmented Generation and KGs for Medical NLP (Target: 10)

- [64][C] Lewis, P., et al. (2020). Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. NeurIPS 2020.
- [65][C] Asai, A., et al. (2023). Self-RAG: Learning to Retrieve, Generate, and Critique through Self-Reflection. arXiv:2310.11511
- [66][C] Wu, Q., et al. (2024). GraphRAG: Unlocking LLM discovery in the private data landscape. Microsoft Research Blog / arXiv.
- [67][V] Xiong, G., et al. (2024). Benchmarking Retrieval-Augmented Generation for Medicine. arXiv:2402.13178
- [68][V] Lu, S., et al. (2024). MedGraphRAG: Graph-RAG for Medical Reasoning. arXiv (verify ID).
- [69][V] Edge, D., et al. (2024). From Local to Global: A Graph RAG Approach to Query-Focused Summarization. arXiv:2404.16130
- [70][V] Gao, Y., et al. (2024). Retrieval-Augmented Generation for Large Language Models: A Survey. arXiv:2312.10997
- [71][V] Belyaeva, A., et al. (2024). Synthesizing Text-to-SQL Data with LLMs for Clinical RAG. Verify venue.
- [72][V] Jin, Q., et al. (2023). Med-PaLM 2: Medical question answering with RAG. Verify venue (Nature or NEJM AI).
- [73][S] Search query: ("RAG" OR "retrieval augmented") AND ("medical" OR "clinical" OR "health") site:arxiv.org - run in Week 1.

4.7 Medical Natural Language Inference (Target: 6)

- [74][C] Tannenbaum, D., et al. (2018). MedNLI - A Natural Language Inference Dataset for the Clinical Domain. Proceedings of ACL 2018 (Clinical NLP Workshop).
- [75][C] Romanov, A., & Shivade, C. (2018). Lessons from Natural Language Inference in the Clinical Domain. arXiv:1808.06752
- [76][C] Velupillai, S., et al. (2018). Adversarial MedNLI: A New Adversarial Dataset for the Clinical Natural Language Inference Domain. Proceedings of the Clinical NLP Workshop at NAACL 2018.

- [77][C] Conneau, A., et al. (2018). XNLI: Evaluating Cross-lingual Sentence Representations (includes Bangla). Proceedings of EMNLP 2018.
- [78][V] Bedi, S., et al. (2024). Multilingual medical NLI: a benchmark. Verify venue.
- [79][S] Search query: "medical NLI" OR "clinical NLI" OR "MedNLI" site:aclanthology.org - run in Week 1.

4.8 Readability Assessment and Control (Target: 6)

- [80][C] Kincaid, J. P., et al. (1975). Derivation of New Readability Formulas for Navy Enlisted Personnel (Flesch-Kincaid). Naval Technical Training Command Report.
- [81][C] Mc Laughlin, G. H. (1969). SMOG Grading - A New Readability Formula. Journal of Reading, 12(8), 639-646.
- [82][C] Xu, W., et al. (2016). Optimizing Statistical Machine Translation for Text Simplification. TACL 4 (introduces SARI).
- [83][V] Deutsch, T., et al. (2020). Linguistic features for readability assessment. Verify venue.
- [84][V] Sinha, M. (2016). A Study of Readability of Texts in Bangla. Education and Information Technologies. Verify DOI.
- [85][S] Search query: "readability" AND ("Bangla" OR "Bengali" OR "low-resource") site:springer.com OR site:sciencedirect.com - run in Week 1.

4.9 General Text Simplification (Target: 6)

- [86][C] Alva-Manchego, F., et al. (2020). Data-Driven Sentence Simplification for Compounding and Splitting. Proceedings of ICNLP 2020.
- [87][C] Saggin, H. (2017). Automatic Text Simplification. Synthesis Lectures on Human Language Technologies, Morgan & Claypool.
- [88][C] Maddela, M., & Xu, W. (2023). A Survey on Text Simplification. Verify venue (likely ACM Computing Surveys or TACL).
- [89][V] Scarton, C., & Specia, L. (2018). A reading comprehension corpus for automatic text simplification evaluation. Verify venue (likely LREC 2018).
- [90][V] Zhang, X., & Lapata, M. (2017). Sentence Simplification with Deep Reinforcement Learning. Proceedings of EMNLP 2017.
- [91][S] Search query: "text simplification" AND ("survey" OR "review") site:aclanthology.org - run in Week 1.

4.10 Low-Resource NLP Methods (Target: 6)

- [92][C] Joshi, P., et al. (2020). The State and Fate of Linguistic Diversity and Inclusion in the NLP World. Proceedings of ACL 2020.
- [93][C] Lauscher, A., et al. (2020). Zero-Shot Cross-Lingual Transfer Improves Zero-Shot Multilingual Task Performance. Proceedings of EMNLP 2020.
- [94][C] Pfeiffer, J., et al. (2020). MAD-X: An Adapters-Based AutoML Framework for Cross-Lingual Transfer. Proceedings of EMNLP 2020 (demo).
- [95][V] Winata, G. I., et al. (2023). Multilingual LLMs: a survey. Verify venue.
- [96][V] Ahmadi, S., et al. (2024). Low-resource NLP: a survey. Verify venue.
- [97][S] Search query: "low-resource" AND ("NLP" OR "language model") AND ("Bangla" OR "Bengali" OR "South Asian") site:arxiv.org - run in Week 1.

4.11 LLM and SSM Backbones (Baselines) (Target: 8)

- [98][C] Xue, L., et al. (2021). mT5: A Massively Multilingual Pre-trained Text-to-Text Transformer. Proceedings of NAACL 2021.
- [99][C] Liu, Y., et al. (2020). Multilingual Denoising Pre-training for Neural Machine Translation (mBART). TACL 8.
- [100][C] Raffel, C., et al. (2020). Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer (T5). JMLR 21.
- [101][C] Devlin, J., et al. (2019). BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. Proceedings of NAACL 2019.

- [102][C] Touvron, H., et al. (2023). Llama 2: Open Foundation and Fine-Tuned Chat Models. arXiv:2307.09288
- [103][C] Llama Team (2024). The Llama 3 Herd of Models. arXiv:2407.21783
- [104][C] Yang, A., et al. (2024). Qwen2.5 Technical Report. arXiv:2412.15115
- [105][C] Google. (2024). Gemma 2: Improving Open Language Models at Practical Sizes. arXiv:2408.00118

4.12 Evaluation Metrics and Human Evaluation (Target: 8)

- [106][C] Lin, C.-Y. (2004). ROUGE: A Package for Automatic Evaluation of Summaries. Text Summarization Branches Out (ACL Workshop 2004).
- [107][C] Papineni, K., et al. (2002). BLEU: A Method for Automatic Evaluation of Machine Translation. Proceedings of ACL 2002.
- [108][C] Zhang, T., et al. (2020). BERTScore: Evaluating Text Generation with BERT. Proceedings of ICLR 2020.
- [109][C] Banerjee, S., & Lavie, A. (2005). METEOR: An Automatic Metric for MT Evaluation. Proceedings of ACL 2005 (Workshop).
- [110][C] Landis, J. R., & Koch, G. G. (1977). The Measurement of Observer Agreement for Categorical Data. Biometrics 33(1), 159-174.
- [111][C] Fleiss, J. L. (1971). Measuring Nominal Scale Agreement Among Many Raters. Psychological Bulletin 76(5), 378-382.
- [112][V] Santhanam, S., et al. (2024). Human evaluation in NLP: a survey. Verify venue.
- [113][S] Search query: "human evaluation" AND ("NLP" OR "text generation") AND ("reliability" OR "agreement")
site:aclanthology.org - run in Week 1.

4.13 Systematic Review Methodology (Target: 8)

- [114][C] Page, M. J., et al. (2021). The PRISMA 2020 Statement. BMJ 372, n71.
- [115][C] Rethlefsen, M. L., et al. (2021). PRISMA-S: An Extension for Reporting Literature Searches. Systematic Reviews 10(1), 39.
- [116][C] Guyatt, G. H., et al. (2008). GRADE: An Emerging Consensus. BMJ 336, 924-926.
- [117][C] Sterne, J. A. C., et al. (2019). RoB 2: A Revised Tool for Assessing Risk of Bias in Randomised Trials. BMJ 366, 14898.
- [118][C] Sterne, J. A. C., et al. (2016). ROBINS-I: A Tool for Non-Randomised Studies. BMJ 355, i4919.
- [119][C] Cooke, A., Smith, D., & Booth, A. (2012). Beyond PICO: The SPIDER Tool. Qualitative Health Research 22(10), 1435-1443.
- [120][C] Thomas, J., & Harden, A. (2008). Methods for the Thematic Synthesis of Qualitative Research in Systematic Reviews. BMC Medical Research Methodology 8, 45.
- [121][C] Shea, B. J., et al. (2017). AMSTAR 2: A Critical Appraisal Tool for Systematic Reviews. BMJ 358, j4008.

4.14 Bibliography Summary

Category	Target	Confirmed [C]	Verify [V]	Search [S]
4.1 Core Clinical TS	18	10	8	0
4.2 Multilingual TS	10	5	4	1
4.3 Bangla NLP	15	6	8	1
4.4 Factuality	12	8	3	1
4.5 Constrained Decoding	8	4	3	1
4.6 RAG / KG	10	2	7	1
4.7 Medical NLI	6	4	1	1

Category	Target	Confirmed [C]	Verify [V]	Search [S]
4.8 Readability	6	3	2	1
4.9 General TS	6	3	2	1
4.10 Low-Resource NLP	6	3	2	1
4.11 LLM Backbones	8	8	0	0
4.12 Evaluation	8	6	1	1
4.13 SR Methods	8	8	0	0
TOTAL	121	70	41	10

Table 4.14 - Bibliography summary. 70 confirmed, 41 to verify, 10 search queries. Total target: 121.

The Week 1-2 search phase (per the review timeline in PDF 1, Section 15) is expected to convert all 41 [V] entries to either [C] (confirmed) or removed, and to execute all 10 [S] search queries to discover additional papers. The final confirmed bibliography is expected to contain 100-120 papers, which exceeds the 100+ target.

5. Import and Verification Workflow

This section specifies the step-by-step workflow for importing papers into Zotero, verifying their bibliographic details, and preparing the library for bibliomatrix generation. The workflow is designed to be executed during Weeks 1-2 of the review timeline.

5.1 Step 1: Bulk Import via Database Export

For each of the 5 primary databases (ACL Anthology, PubMed, IEEE Xplore, ACM DL, Scopus), execute the search string from PDF 1 Section 5, then export the results as RIS or BibTeX. Import each file into Zotero via File -> Import. Zotero automatically creates a collection per imported file; move items into the appropriate collection from Section 2.1.

5.2 Step 2: DOI Bulk Verification

Python script: `verify_dois.py` (run after Zotero import)


```
#!/usr/bin/env python3
"""Bulk-verify DOIs for all Zotero items via Crossref API."""
import requests, json, time
from pyzotero import zotero

ZOTERO_LIBRARY_ID = "YOUR_LIBRARY_ID"
ZOTERO_API_KEY = "YOUR_API_KEY"
zot = zotero.Zotero(ZOTERO_LIBRARY_ID, "user", ZOTERO_API_KEY)

items = zot.everything(zot.items())
for item in items:
    doi = item["data"].get("DOI", "")
    if not doi:
        print(f"MISSING DOI: {item['data'].get('title', '')[:60]}")
        continue
    r = requests.get(f"https://api.crossref.org/works/{doi}",
                    headers={"User-Agent": "ReviewBot/1.0 (mailto:you@x.com)"})
    if r.status_code == 200:
        crossref = r.json()["message"]
        print(f"OK: {doi} -> {crossref['title'][0][:60]}")
    else:
        print(f"FAIL: {doi} (status {r.status_code}) -> {item['data'].get('title', '')[:60]}")
    time.sleep(1) # respect Crossref rate limit

# Output: verify_dois_log.txt with OK/FAIL/MISSING per item
# Move all FAIL and MISSING items to 00_To_Verify collection
```

5.3 Step 3: PDF Attachment and OCR

For each verified item, download the PDF (open access via ACL Anthology, arXiv, or PMC; subscription via institutional access) and attach it to the Zotero record. ZotFile automatically renames the PDF to the citation key format (authorlastname_year_shorttitle.pdf). For scanned PDFs without a text layer, run OCR via Zotero's built-in OCR (or Tesseract externally) before attachment so the PDF is searchable.

5.4 Step 4: Tagging and Collection Assignment

For each verified item with attached PDF, apply: (i) one primary collection from Section 2.1; (ii) at least one topic tag from Section 2.2; (iii) optional secondary collections for cross-cutting relevance. Use Zotero's batch-select feature (Ctrl+click multiple items, then drag to a collection) to speed this up. The target is 100% of items tagged with at least one collection and one tag by end of Week 2.

5.5 Step 5: Duplicate Detection and Resolution

Run Zotero's duplicate detection (Edit -> Find Duplicates) after all imports. Common duplicates arise from importing the same paper from multiple databases. Merge duplicates carefully: keep the record with the most complete metadata (usually the Crossref-verified one) and the best PDF. After merging, re-verify tags and collections on the merged record.

5.6 Step 6: Better BibTeX Export for Bibliometrix

Once the library is clean (all items verified, tagged, and deduplicated), export the entire library as BibTeX via Better BibTeX: File -> Export Library, choose Better BibTeX format, save as review_refs.bib. This file is the input to the R bibliometrix workflow in Section 3.5. Re-export after any library change to keep the bibliography in sync with the bibliometrix.

6. Bibliomatrix Generation Checklist

This is the concrete checklist for generating the four bibliomatrices during Week 3 (synthesis phase) of the review timeline. Each matrix has a defined input, tool, output format, and target location in the review paper.

#	Matrix	Input	Tool	Output	Paper Location
M 1	Paper x Factuality Mechanism	Zotero items + manual classification	Manual in Google Sheets (R script for auto-extraction)	Heat-map table	Section 4 (Results - Methods Landscape)
M 2	Paper x Bangla Relevance	Zotero items + manual classification	Manual in Google Sheets	Stacked bar chart	Section 5 (Results - Resources Landscape)
M 3	Co-Citation Network	review_refs.bib (Better BibTeX export)	VOSviewer (co-citation analysis)	Network PNG + PDF	Section 3 (Results - Flow) or Appendix
M 4	Keyword Co-Occurrence	review_refs.bib (author keywords)	VOSviewer (co-occurrence analysis)	Network PNG + PDF	Section 4 (Results - Methods Landscape)
M 5	Annual Scientific Production	review_refs.bib	Bibliometrix R (annual ScientificProduction)	Line chart PNG	Section 3 (Results - Flow)
M 6	Three-Field Plot (AU x DE x SO)	review_refs.bib	Bibliometrix R (threeFieldsPlot)	Sankey diagram PNG	Section 4 (Results - Methods Landscape)

Table 6.1 - Six bibliomatrix outputs with their tools and target locations in the review paper.

6.1 Quality Checks for Bibliomatrices

- **Completeness:** Every included paper appears in at least one matrix.
- **Consistency:** The same paper is classified the same way across all matrices (e.g., a paper tagged "factuality" in Zotero is counted in the factuality column of M1).
- **Reproducibility:** All R scripts and VOSviewer config files are committed to the review repo with the BibTeX input, so any reviewer can regenerate every matrix.
- **Visual quality:** All matrices exported at 600 DPI for PNG and vector PDF; font sizes legible at 100% zoom in the printed review paper.
- **Caption:** Every matrix has a numbered caption (Figure 1, Figure 2, ...) and is referenced in the body text.

7. Integration with Review Timeline

This bibliography workstream runs in parallel with the main review timeline specified in PDF 1 Section 15. The mapping below shows which bibliography tasks fall in which week.

Week	Main Review Task (from PDF 1)	Bibliography Task (this guide)
W1	Finalize RQs, build search strings, submit PROSPERO	Set up Zotero library (11 collections, 18 tags, plugins); bulk import from 5 databases
W2	Execute searches, dedup, title/abstract screening	DOI bulk verification (verify_dois.py); PDF attachment; tagging; deduplication; export review_refs.bib
W3	Full-text review, data extraction, RoB assessment	Generate M1-M6 bibliomatrices; manual classification for M1 and M2; integrate into draft
W4	SWiM synthesis, GRADE rating, write first draft	Finalize matrix captions; commit R scripts and BibTeX to repo; verify all 100+ references are cited

Table 7.1 - Integration of bibliography tasks with the 4-week review timeline.

8. Deliverables Checklist

By the end of Week 2, the following deliverables must be complete:

- **D1.** Zotero library with 11 collections, 18 controlled-vocabulary tags, and 100+ verified items (all [C] status; no [V] remaining).
- **D2.** Every item has an attached PDF (searchable text layer; OCR if needed).
- **D3.** Every item is assigned to at least one primary collection and one topic tag.
- **D4.** Duplicate detection run; all duplicates merged.
- **D5.** review_refs.bib exported via Better BibTeX; committed to the review repo.
- **D6.** verify_dois_log.txt committed; all FAIL/MISSING items resolved or moved to 00_To_Verify with a note.
- **D7.** Zotero library backed up to private GitHub repo (zotero.sqlite + storage/).

By the end of Week 3, the following additional deliverables must be complete:

- **D8.** Matrix M1 (paper x factuality mechanism) populated for all included papers.
- **D9.** Matrix M2 (paper x Bangla relevance) populated for all included papers.
- **D10.** Matrices M3-M6 generated via VOSviewer and Bibliometrix R.
- **D11.** All matrix PNG and PDF files committed to the review repo at 600 DPI.
- **D12.** R scripts (generate_bibliomatrix.R) and VOSviewer config files committed for reproducibility.

9. Closing Note

This guide operationalizes the bibliography management for the 1-Month Systematic Review. The 121-paper target bibliography (Section 4) is the gathering list; the Zotero architecture (Section 2) is the organization system; and the bibliomatrix specification (Section 3) is the analytical output. Together they ensure the review is reproducible, transparent, and defensible to Q1 journal reviewers. All scripts (verify_dois.py, generate_bibliomatrix.R) and the BibTeX export are version-controlled alongside the manuscript so that any reviewer can regenerate every matrix and verify every citation from the public repository.